

Delta Lake



What is Delta Lake?



What is Delta Lake?

What is Delta Lake?

<https://delta.io>

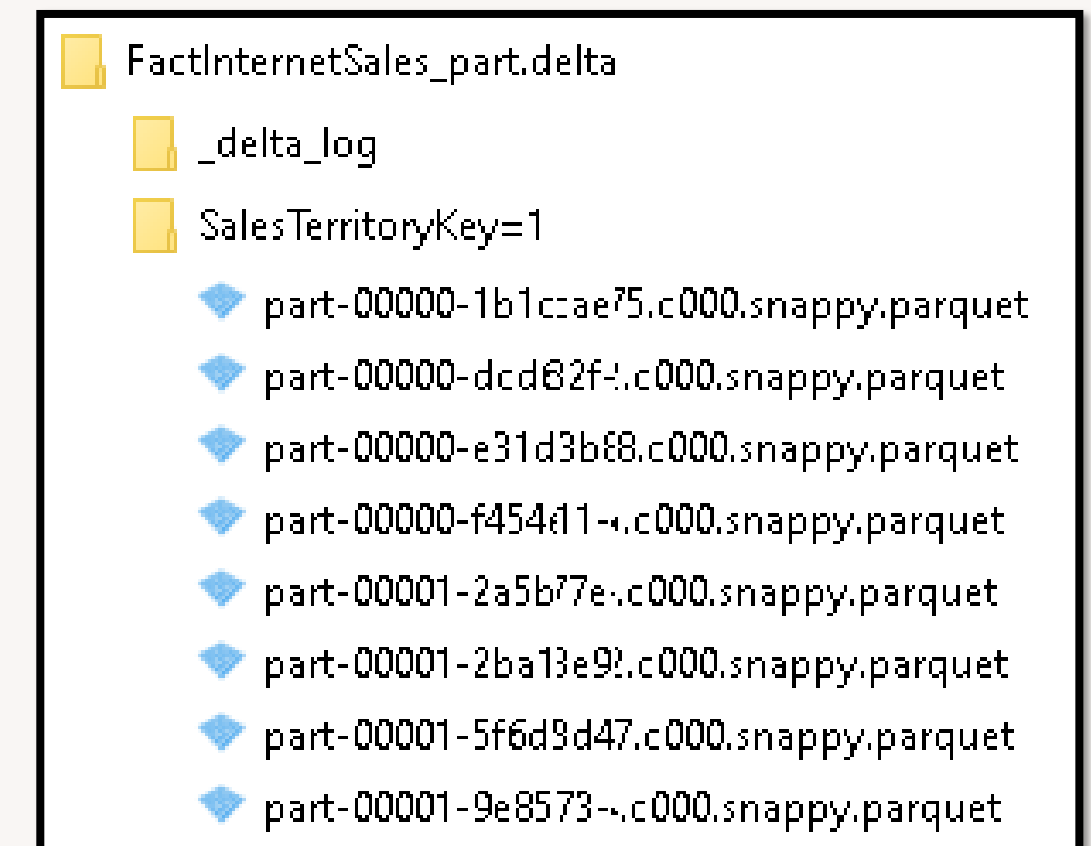
Delta Lake is an open-source storage framework that enables building a Lakehouse architecture with compute engines including Spark, PrestoDB, Flink, Trino, and Hive and APIs for Scala, Java, Rust, Ruby, and Python.

- ACID compliant transactions
 - Optimistic Concurrency Control
- Support for UPDATE / MERGE
- Time-Travel
- Schema enforcement and evolution
- Batch & Streaming
- 100% compatible with Spark

What is Delta Lake?

<https://delta.io>

- Built for Cloud Object stores
- Everything is stored in one folder
 - Meta-data
 - Transaction log / **Delta Log**
 - Data
- Could basically Copy & Paste whole Delta table
- Supports any storage sub-system
- Consumer only needs location



Delta Lake is the default for all
tables created in Databricks

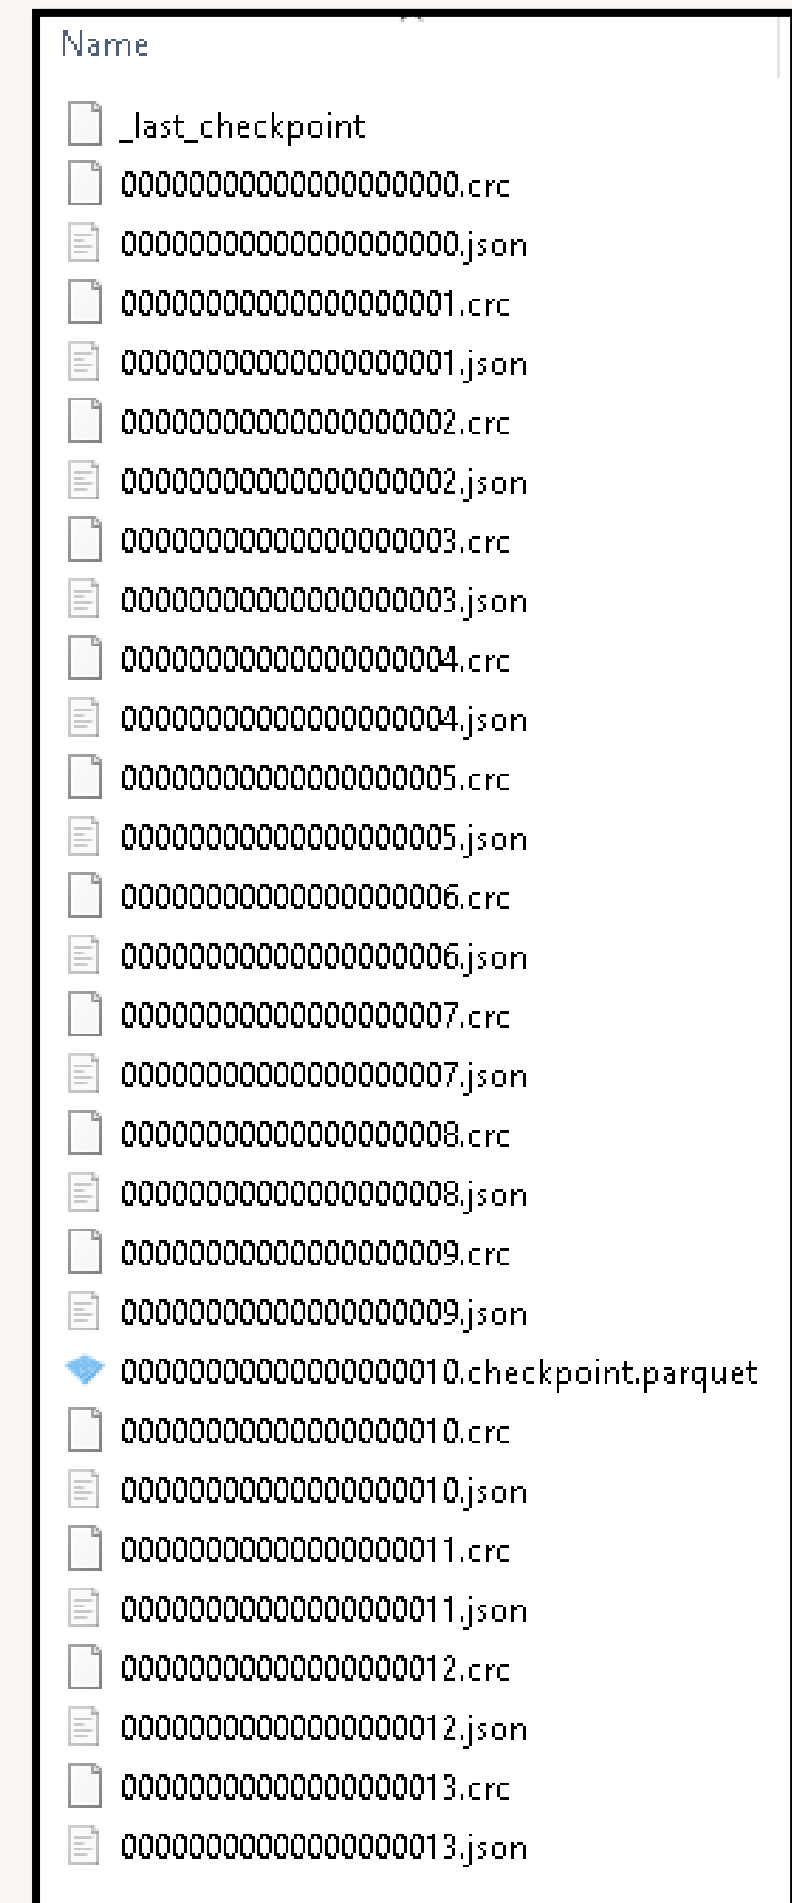
What is the Delta Log?



What is the Delta Log?

The Transactional Layer

- Contains
 - Table schema + changes
 - References to files
 - Metadata and metrics
- Stored as JSON and Parquet
- One file/version per transaction
- Allows [optimistic] concurrency control
- Used for time-travel, streaming, ...



What is the Delta Log?

DESCRIBE HISTORY

 DESCRIBE HISTORY gold.my_big_table

▶ (1) Spark Jobs

Table

Data Profile

	version ▲	timestamp ▲	userId ▲	userName ▲	operation ▲	operationParameters
1	1185	2022-06-13T16:45:39.000+0000			OPTIMIZE	▶ {"predicate": "[]", "zOrderBy": "[]", "batchId": "0", "auto": "false"}
2	1184	2022-06-13T16:18:24.000+0000			VACUUM END	▶ {"status": "COMPLETED"}
3	1183	2022-06-13T16:18:19.000+0000			VACUUM START	▶ {"retentionCheckEnabled": "false", "defaultRetentionMillis": "259200000"}
4	1182	2022-06-13T13:59:19.000+0000			MERGE	▶ {"predicate": "((target.purchase_id = updates.purchase_id) AND (target.iteration < updates.iteration)) OR ((NOT (updates.store_fee_rate = target.store_fee_rate)) OR (NOT (updates.store_fee_description = target.store_fee_description))) OR ((NOT (updates.store_fee_absolute = target.store_fee_absolute)) OR (NOT (updates.net_sales_after_fees = target.net_sales_after_fees))))", "actionType": "update"}, [{"actionType": "insert"}]
5	1181	2022-06-13T13:56:44.000+0000			MERGE	▶ {"predicate": "(target.purchase_id = updates.purchase_id)", "matchedPredicates": {}, "notMatchedPredicates": "[]"}]



How does Delta Lake work?



DML Operations – UPDATE

Product	Price
Notebook	900 €
PC	1,500 €
Tablet	500 €

```
UPDATE DimProduct
SET Price = 1300
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

User

_delta_log


Storage

0000.json
"add": { "part-01.parquet", ... }


0001.json
"remove": { "path": "part-01.parquet", ... },
"add": { "path": "part-02.parquet", ... }



Parquet
part-01
(3 rows)



Parquet
part-01
(3 rows)



Parquet
part-02
(3 rows)

DML Operations – DELETE

User

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

```
DELETE FROM DimProduct
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
Tablet	500 €

_delta_log

0000.json



0001.json

```
"remove": { "path": "part-01.parquet", ... },
"add": { "path": "part-02.parquet", ... }
```

0002.json

```
"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }
```

Storage



DML Operations – INSERT

Product	Price
Notebook	900 €
Tablet	500 €

```
INSERT INTO DimProduct
VALUES ('Monitor', 200)
```

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

User

_delta_log

0000.json
0001.json


0002.json


"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }


0003.json


"add": { "path": "part-04.parquet", ... }


Storage



Parquet
part-01
(3 rows)



Parquet
part-02
(3 rows)


Parquet
part-03
(2 rows)


Parquet
part-01
(3 rows)


Parquet
part-02
(3 rows)


Parquet
part-03
(2 rows)


Parquet
part-04
(1 row)



DML Operations

- Operations are logged in `_delta_log`
 - Old files are **logically(!)** removed
 - New files are added
- Most operations create new files! Even a **DELETE** can!

Can create A LOT of files!

File & Storage Management



Data Management – VACUUM

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

VACUUM DimProduct

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

User

_delta_log

0000.json

⋮
0003.json

0004.json

"add": { "path": "part-04.parquet", ... }



0005.json

{"VACUUM START ", ... "numFilesToDelete": 2, ... }

0006.json

{"VACUUM END ", ... "numDeletedFiles": 2, ... }

Storage



Data Management – OPTIMIZE

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

OPTIMIZE DimProduct

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

User

_delta_log

0000.json

...

0005.json

0006.json

{"VACUUM END ", ... "numDeletedFiles": 2, ... }

0007.json

"remove": { "path": "part-03.parquet", ... }
"remove": { "path": "part-04.parquet", ... }
"add": { "path": "part-05.parquet", ... }

Storage



Parquet
part-03
(2 rows)



Parquet
part-04
(1 row)



Parquet
part-03
(2 rows)



Parquet
part-04
(1 row)



Parquet
part-05
(3 rows)



Data Management

VACUUM

- Physically removes unreferenced files older than X days
- Never touches files of latest version of Delta table!

```
VACUUM events  
[RETAIN num HOURS]  
[DRY RUN]
```

OPTIMIZE

- Collapse small files into bigger files
- Clustering / Ordering
- Improve query performance

```
OPTIMIZE events  
[WHERE date = 20200101]  
[ZORDER BY (eventType)]
```


Data Management

VACUUM and OPTIMIZE

- VACUUM DRY RUN
 - Only shows first 1000 files to be deleted
 - Use SCALA to get the actual number of files to be removed!
- Can take a long time!
- OPTIMIZE
 - works per partition level

```
1 %scala
2 spark.sql("VACUUM gold.my_big_table DRY RUN")
```

▶ (12) Spark Jobs

▶  res2: org.apache.spark.sql.DataFrame = [path: string]

Found 5888 files and directories in a total of 18531 directories to be removed
res2: org.apache.spark.sql.DataFrame = [path: string]

Command took 1.85 minutes -- by gbrueckl@paiqo.com at 13/06/2022, 20:46:10 CEST

Data Management

RESTORE

- Restores a previous state of the Delta table
- At **version** or **timestamp**
- Meta-data only operation
- Creates a new version

```
RESTORE events  
TO TIMESTAMP AS OF  
'2022-05-03'
```

CLONE

- **SHALLOW** or **DEEP**
- Forks Delta Log
 - **DEEP**: copies data files
 - **SHALLOW**: references data files
- Ideal for testing

```
CREATE TABLE  
events_clone  
SHALLOW CLONE events;
```

Data Management

RESTORE and CLONE

- You can RESTORE as often as you want
 - To rollback another RESTORE
- RESTORE does not create any new [data] files
- DEEP Clones are incremental and can be used for Backup



Demo



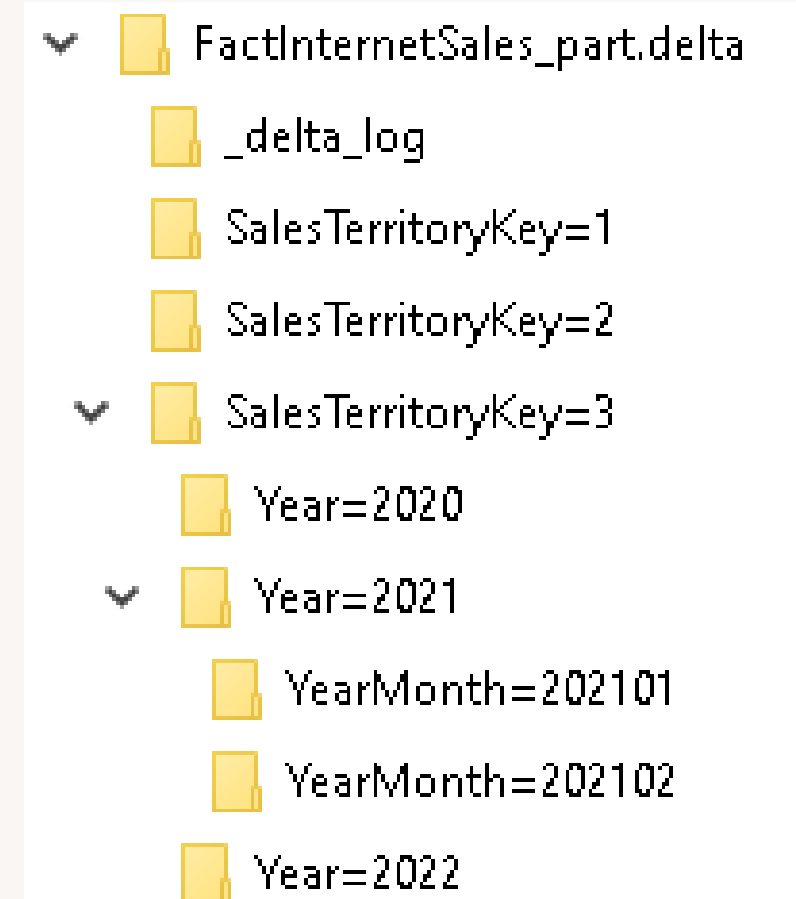
Partitioning



Partitioning

Basics

- Delta Tables can be partitioned
 - For ETL performance (usually on Bronze, Silver)
 - For query performance (usually on Gold)
- **Goal:** touch as few partitions as possible/necessary
 - ETL and Query performance can conflict
 - Explicitly specify Partitioning columns
- Partition by Time [and ?]



Partitioning

Advanced

- Avoid over-partitioning!
 - < few 1000s partitions
 - Single partition should be > 1 GB
- Use generated columns
 - `EventTimestamp` → partition by `CAST(EventTimestamp AS DATE)`
 - Delta engine will [try to] push filters on `EventTimestamp` down to partition
- Used to separate transactions and processing jobs
 - Explicitly specify partitions you touch (e.g. `MERGE` target)!
 - Check Delta Log history for query predicates!

Partitioning

Advanced

- Physical .parquet file does not contain the partitioning-columns!
- **path** could point anywhere!
- You do not need to specify all partitioning columns sequentially!

_delta_log Entry

```
{
  "add": {
    "path": "SalesTerritoryKey=8/SalesDate=20220103/part-....",
    "partitionValues": {
      "SalesTerritoryKey": "8",
      "SalesDate": 20220103
    },
    "size": 114365,
    "modificationTime": 1611740902000,
    "dataChange": true,
  }
}
```

Streaming



Streaming

Basics

- Delta Lake can be used as source and target for streaming
- It's technically still [micro-]batches
 - As is Spark Streaming
- Streaming works on a file-level
- Files are processed in order of
 - Version/Transaction number
 - File index (part-XXXXX...snappy.parquet)

Streaming

Advanced

- Checkpoints
 - Track state of what has already been processed from source
 - One checkpoint per source
 - Could stream from same source multiple times using different checkpoints
- **MERGE** only with **foreachBatch()**
- Control the Trigger/Batch size!
- Avoid **Trigger.Once**
- Can stop/resume stream at any time

Delta Lake TableProperties



Delta Lake Table Properties

- Can be defined on different levels
- Table Properties
 - `delta.autoOptimize.optimizeWrite`
 - `spark.databricks.delta.properties.defaults.optimizeWrite`
(default for new tables)
- Configured Settings during Execution
 - `spark.databricks.delta.optimizeWrite.enabled`
- Execution settings have priority over table properties!

Delta Lake Table Properties

Important Table Properties to know

- `delta.appendOnly`
- `delta.autoOptimize.autoCompact`
- `delta.autoOptimize.optimizeWrite`
- `delta.deletedFileRetentionDuration`
- `delta.logRetentionDuration`
- `delta.dataSkippingNumIndexedCol`

Delta Lake Table Properties

- Use defaults for commands
 - Define exceptions on table level
- No need to use individual commands per table
- Changing table properties are also a Delta transaction

Conclusion

Take Aways

- Delta Lake can solve a lot of problems for you
- File management is crucial
- Data maintenance jobs are mandatory
- Use table properties